


CLOUDLET'S SERVERLESS

From Knative to the Portal

What we built, and why.

Team review · September 2026



TODAY

What we'll talk about

- What is serverless computing
- What is Knative and OpenShift Serverless Operator
- The API: active-active Knative wrapper
- Functions: buildpacks, kpack, build controller
- The portal and cloudlet-apis
- Where we go next

40

slides · six topics · one platform

Serverless, defined

- The servers are still there — you stop owning them.
- You hand over code. The platform runs it only while work arrives.

serverless $\neq$ no servers

someone else's problem, on purpose

What that buys you

- Scale to zero when idle
- Scale out when traffic arrives
- Pay for requests, not hours
- Nothing to patch, nothing to plan
- One deploy, no capacity meeting

0 → N → 0

replicas follow traffic

Serverless: the trade-offs

	PRO	CON
Cost	Pay only when it runs	Bills spike under load
Ops	No servers to patch	Runtime is not yours
Scale	Up and down by itself	Cold start on first hit
Speed	A function in minutes	Platform lock-in
Fit	Bursty, event-driven	Long-running jobs, state
Deploy	Easy, few clicks	Less control and customization



Knative, and where it came from

- Google open-sourced Knative in 2018; it joined the CNCF in 2022.
- It is not a platform of its own — it is serverless building blocks for Kubernetes.
- Serving handles requests, Eventing routes events. We run Serving only.

2018

Google open-sources Knative

2022

donated to the CNCF

Serving

request-driven workloads

Eventing

brokers, triggers, CloudEvents

Ours

Serving, via the Operator


Knative, from what we already run

- We know Deployment, Service, Ingress, HPA
- Knative bundles them into one object
- Adds revisions, routing and scale-to-zero
- Runs on any Kubernetes, ours is OpenShift

NEED	PLAIN KUBERNETES	KNATIVE
Run the code	Deployment	Revision
Reach it	Service + Ingress	Route
Scale it	HPA, never to zero	KPA, down to zero
Change it	Rolling update	New Revision, old kept
Own hostname	Ingress rule	DomainMapping

One object: the Knative Service

- The Service owns Configuration and Route
- Configuration holds the desired code
- Every change is an immutable Revision
- Route sends traffic to revisions
- DomainMapping gives it your hostname



The Service, in one file

```
apiVersion: serving.knative.dev/v1
kind: Service
metadata:
  name: billing
  namespace: team-a-serverless          # one namespace per SSO group
spec:
  template:
    metadata:
      annotations:
        autoscaling.knative.dev/metric: concurrency
        autoscaling.knative.dev/target: "80"
        autoscaling.knative.dev/min-scale: "0"      # scale to zero
        autoscaling.knative.dev/max-scale: "10"
    spec:
      containers:
        - image: registry.central/team-a/billing@sha256:9f2c...
          ports:
            - containerPort: 8080
          resources:
            requests: { cpu: 250m, memory: 512Mi }
```


The pieces that do the scaling

- Concurrency: requests in flight per pod
- The queue proxy measures it
- The KPA turns that into replicas
- The Activator stands in at zero
- minScale and maxScale bound it

PIECE	WHAT IT DOES	ROLE
Queue proxy	Counts in-flight requests	the meter
KPA	Requests → replica count	the decision
Activator	Holds requests while pods start	the safety net
minScale	Keeps N pods warm	the override
maxScale	Caps the fan-out	the ceiling

Scaling to zero, and back

- An idle revision sits at zero pods
- A request arrives for that revision
- The Activator holds it, never drops it
- The KPA is asked for capacity
- Pods start, the request is served
- Once healthy, traffic goes direct



The OpenShift Serverless Operator

- Red Hat's supported Knative distribution
- One KnativeServing custom resource
- OLM installs it and upgrades it
- Kourier ingress, Routes created for you
- Catalog mirrors with oc-mirror

Knative Serving

Service · Revision · Route ·
DomainMapping

KnativeServing CR

one document

Serverless Operator

reconciles, upgrades

Operator Lifecycle Manager

Subscription → catalog

OpenShift

one cluster per region

Why the Operator

UPSTREAM KNATIVE

SERVERLESS OPERATOR

Upgrades

YAML we maintain by hand

OLM, from the catalog

Ingress

We run and tune it

Kourier, managed

Routes, TLS

Hand-made, extra RBAC

Created for us

Airgap

Image by image

oc-mirror, one path

Support

Community

Red Hat

Who owns it

Us

The platform

The API

One HTTP call, a serverless application.

What the API is

- One HTTP API in front of Knative
- Two offerings: functions and containers
- It validates, then answers 202
- One call reaches both regions
- Nobody needs cluster access

one API

FastAPI, running in both regions

functions

your git repo – we build it

containers

your image – we run it

the contract

202 + statusUrl, closed statuses

the reach

central + south, one call

The endpoint structure

- 23 endpoints under one base path
- The group is in the path
- Both offerings have the same shape
- OIDC token or admin key
- Allow live streaming with short-lived tickets

```
/api/serverless/v1
/groups/{group}/functions
  POST · GET · PUT · DELETE
  /{name}/build
  /{name}/stats · /pods · /logs
/groups/{group}/containers
  POST · GET · PUT · DELETE
  /{name}/pull
  /{name}/stats · /pods · /logs
/functions/info · /containers/info
/stream-tickets
```


Anatomy of a workload

- Identity: name, hostname, regions
- Source: git repo or image
- Runtime: env, files, port, size
- Scaling: min, max, metric, target
- Response: same shape, secrets redacted

Identity`name · hostname · regions[]`**Source**`gitRepo · branch · runtime | image`**Runtime**`env[] · files[] · port · size`**Scaling**`minScale · maxScale · metric · target`**Response**`the request, secrets redacted`

Every workload has one status

- A closed set, published on /info
- Pending, Building, Deploying, Ready
- Building only exists for functions
- Failed always carries a reason
- One row per region, rolled up



You send this

POST /api/serverless/v1/groups/team-a/functions

Authorization: Bearer <keycloak token>

```
{
  "name": "billing", "port": 8080, "size": "small",
  "gitRepo": "https://git.internal/team-a/billing.git",
  "branch": "main", "runtime": "python", "version": "3.12",
  "regions": ["central", "south"],
  "env": [{ "name": "LOG_LEVEL", "value": "info" }],
  "scaling": { "minScale": 0, "maxScale": 10,
               "metric": "concurrency", "target": 80 }
}
```

no namespace, no image, no Knative, no kubectl

You get this back, immediately

202 Accepted

```
{  
  "name": "billing", "status": "Pending",  
  "hostname": "billing-team-a.serverless.example",  
  "statusUrl": "/api/serverless/v1/groups/team-a/functions/billing",  
  "regions": [{ "region": "central", "status": "Pending" },  
               { "region": "south", "status": "Pending" }]  
}
```

validated synchronously – a 400 here means nothing was written

accepted, not finished: the work happens after the answer

Then you poll the same URL

GET /api/serverless/v1/groups/team-a/functions/billing


200 OK

```
{
  "name": "billing", "status": "Ready",
  "url": "https://billing-team-a.serverless.example",
  "image": "registry.central/team-a/billing@sha256:9f2c...",
  "scaling": { "minScale": 0, "maxScale": 10 },
  "regions": [
    { "region": "central", "status": "Ready", "revision": "billing-00007" },
    { "region": "south", "status": "Ready", "revision": "billing-00007" }],
  "lastBuild": { "status": "Succeeded", "finishedAt": "..." }
}
```

Pending → Building → Deploying → Ready, same shape every time


Talking to the cluster

- Client certificate mTLS, always
- No kubeconfig, no service-account token
- The certificate's CN is the user
- Address derived from the region name
- Every write is server-side apply




Active/active, two regions

- The API runs in both regions
- One deploy fans out to both
- No leader election, deterministic specs
- A region builds what it runs
- State lives in the Knative Service




The tenant controller

- A namespace per SSO group
- Separate process, separate certificate
- The API never creates namespaces
- Provision before every deploy, fail closed
- Reconcile the local cluster only
- Garbage-collect empty namespaces, opt-in



Backups: Trident Protect

- Trident Protect protects whole namespaces
- One Application object per tenant namespace
- A Schedule runs it unattended
- A Snapshot stays on the storage
- A Backup leaves for object storage
- Restore puts the namespace back




Functions & builds

'Here is my repo' hides a build.

Why a function needs a build

- A cluster cannot run a repository
- Kubernetes only ever starts images
- Someone must turn source into an image
- That step is the build
- For containers, the user already did it





Buildpacks

- An open standard, not a Dockerfile
- Detection reads your source
- The same base image for everyone
- Rebase swaps that base, no rebuild
- A CNCF project, like Knative

	DOCKERFILE	BUILDPACKS
Who writes it	Every team, again	Nobody
Base image	Whatever was pinned	One curated stack
Language setup	Copied between repos	Detected from source
Patching the base	Rebuild and hope	Rebase, in seconds
Consistency	Per repository	Platform-wide

Paketo: the buildpacks we run

- The open-source implementation we use
- One buildpack per language, plus shared ones
- Detection picks the group for you
- We mirror them into the airgap
- Python, Go and Node today

BUILDPACK	DETECTS	GIVES YOU
paketo/python	requirements.txt	CPython, pip install
paketo/go	go.mod	a compiled binary
paketo/nodejs	package.json	Node, npm ci
paketo/ca-certificates	always	our internal CA
paketo/procfile	Procfile	the start command

kpac: buildpacks on Kubernetes

- Buildpacks need something to run them
- kpac runs them as pods, in-cluster
- You declare an Image, it builds
- It rebuilds on commit or base change
- The output is a digest, nothing else

kpac

a Kubernetes controller

declarative

you write an Image, not a pipeline

in-cluster

one pod per build


rebuilds

new commit · new base · new buildpack

output


a registry digest

The kpack resources



A build, phase by phase

- The lifecycle runs as init containers
- Each phase is a named container
- Each phase has its own log
- Cached layers make a rerun fast
- Export assemblies and pushes the image



Our kpack chart

- Not a fork, a Helm chart
- Packages upstream kpack faithfully
- CRDs as templates, for the webhook
- Creates the cluster build content
- Mirror scripts for the airgap

the chart

upstream kpack 0.18

CRDs

templated, not an untouched crds/

clusterBuild

ClusterStack + ClusterStore

credentials


projected by external-secrets

scripts/mirror

images, buildpackages, runtimes


How the API uses kpack

- The API declares, it never builds
- One Image per function, per region
- A git Secret for the clone
- A build ServiceAccount for the registry
- Owned by the Knative Service



The build controller

- Watches Image.status.latestImage
- Applies the digest onto the Service
- Both ends stay in one region
- Only ever writes digests
- Ships with no web stack




Portal & the future

A console, and a library for the next API.


How the portal wires it

- Next.js 16, server components
- Keycloak token stays server-side
- The SSO group is the project
- Forms are built from /info
- Live logs over SSE tickets
- Airgap-native: no CDN, no fonts



A console, not a Serverless UI

- Serverless is live, Compute is next
- An offering is data, not code
- The shell is already generic
- Next: shared layer, generated client



cloudlet-apis: start at step five

- Extracted from the Serverless API
- core, web and auth extras
- Layering enforced by tests
- A new API wires in five steps
- Same group rules everywhere

[auth]

OIDC · admin key · token proxy ·
tickets

[web]

health · error envelope · offline docs

core

errors · names · logging · request id

Code in, a running service out.

4

repositories

3

services

23

endpoints

2

regions

3

runtimes

0

kubectl

- Operator owns Knative; we own the contract
- Deterministic specs, so no leader
- Builds declared, watched, rolled in
- The second product is cheaper

**Thank you.
Questions?**